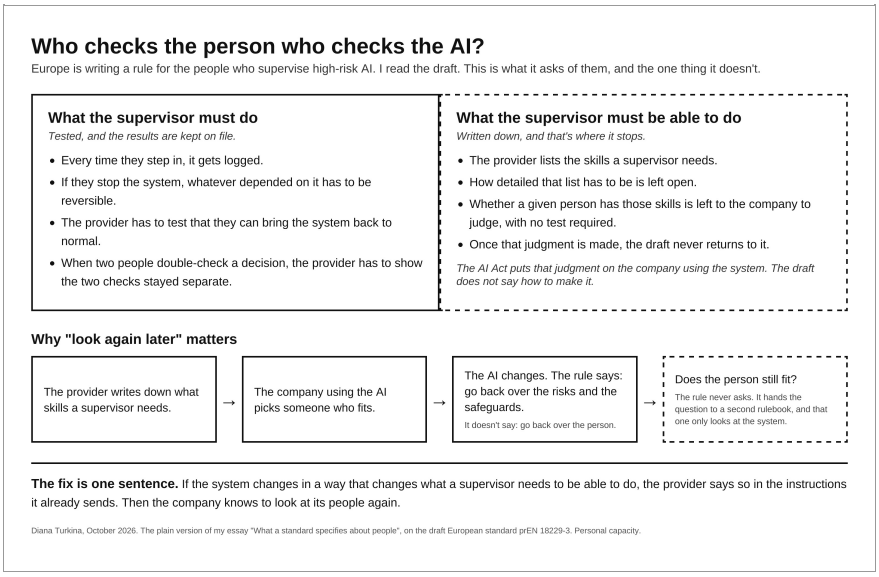
What a standard specifies about people

A close reading of prEN 18229-3 (human oversight) and prEN 18228 (AI risk management), from sixteen public-enquiry comments

Diana Turkina
12 October 2026 · Version 1.0

The European draft standard on human oversight is careful about what the overseer does. It logs their interventions, confirms their actions, rolls back what they undo. It is much less careful about whether they can do the job, and it never asks whether a judgment that they can is still true a year later. I read it during the enquiry and this is what I found.



Abstract. Draft EU standard prEN 18229-3 specifies what a human overseer of a high-risk AI system must do. It specifies little about what the overseer must be able to do, and nothing about when a finding that they can stops holding. I read the draft clause by clause during its public enquiry and wrote sixteen comments. Two exhibits stand out. In 5.8, a provider must file test results showing that two manual reviewers used separate accounts, but need only define the reviewers' competencies. In 5.2.6, a review of oversight assignments is triggered by system changes but does not cover the person assigned, and it refers to prEN 18228 Clause 11, whose inputs are system-side only. The gap falls at the deployer's designated person, whom no provider quality system reaches. I propose one sentence for each document. Four of the comments are reproduced in full. This is a record of observations made in a personal capacity on one enquiry draft. It does not report the outcome of the working group's comment resolution.

Keywords. human oversight; AI Act Article 14; harmonised standards; competence assurance; competence reassessment; deployer; prEN 18229-3; prEN 18228.

Provenance. I read prEN 18229-3 and prEN 18228 clause by clause during the public enquiry on the first, which closed on 22 September 2026, and wrote sixteen comments in the enquiry format, which I submitted. Comment resolution happens inside the working group and is not visible outside it, so this document records what the drafts said at that date, which clauses I read them against, and what I proposed.

Quotation practice. Draft European standards are licensed documents, free to read only during the enquiry window. This essay cites clauses by number and title, paraphrases their requirements, and quotes at most a few words at a time, for the purpose of comment and analysis. The comments themselves are my own text and four of them are reproduced in the Appendix.

Capacity. Written in a personal capacity. No employer affiliation or endorsement. This document is not the position of any standards body or committee. The analysis builds on a five-framework study deposited at SSRN (abstract 7248205) with its coding sheet and term-search log at doi.org/10.5281/zenodo.22075225, CC BY 4.0.

Cite as. Turkina, D. (2026). What a standard specifies about people. Version 1.0. Zenodo, doi.org/10.5281/zenodo.23311590.

1 · The draft is precise about people

prEN 18229-3 is the part of the European AI standards family that covers human oversight. Drafted under the standardisation request supporting the AI Act, it will in time tell providers of high-risk systems what Article 14 requires of them in practice. This essay is about a pattern visible only when the draft's provisions on people are read together.

The draft is precise about human beings. Much commentary on AI governance assumes documents of this kind gesture at human oversight and move on. This one specifies what the designated person, the deployer's staff member to whom oversight of the system is assigned, does, in detail, with the same apparatus it uses for the system.

At 5.7.2, 5.7.3 and 5.7.4 it requires that interventions by the designated person be logged, that certain actions be confirmed, and that dependent processes be rolled back. Those are requirements about a person's actions, written to the standard applied to machine behaviour: stated, testable, documented in the technical file. At 5.7.5 d) it requires the provider to test that the procedure for returning to normal operation is executable by the designated person. At 5.3.2.3.1 and 5.3.4.4 it requires verification that designated persons can identify hazardous outputs and perform key oversight tasks, with acceptance criteria and results documented in the technical file (6.2 e) and f)). Those are capacity tests, the only two in the draft with a pass condition, and both leave undetermined which tasks are verified, on whom, and under what conditions (my ninth comment, on 5.3.4.4). The draft treats the human as part of the system being specified.

Beyond those two tests it states no requirement of what the person has to be able to do, and nothing in the draft asks whether a finding of competence still holds.

2 · Two exhibits

2.1 Account separation is tested. Competence is declared.

The clearest instance sits within one subclause tree, 5.8, which governs manual review in biometric identification: a person confirms or rejects a match before it takes effect. Under Article 43(1) a provider that applies

the harmonised standard in full may assess conformity itself; a notified body is required only where it does not. The standard's content therefore decides what is checked.

5.8.3, titled *Manual reviewer independence mechanisms*, specifies independence as four measures: individual credentials for each reviewer, prevention of a single account completing both reviews, blinding of the second reviewer to the first result, and separate recording of each review action. Those four measures establish that two separate accounts performed the reviews. That is what independence means in this subclause, and it is narrower than the word. 5.8.4, titled *Documentation requirements*, then requires the technical documentation to carry the implemented mechanisms and test results demonstrating that they function as specified.

5.8.5, a few pages on, specifies what the provider states about manual reviewers. Item a) lists three competencies. They are to be defined and documented. Not tested. Not evidenced. Defined.

So within one document and one subclause tree, the draft treats one property of a reviewer as something to be demonstrated and another as something to be stated. Account separation gets test results. Competence gets a definition. Both bear on whether the second review is worth anything, and only one of them has to be shown.

A declaration can be checked for presence and consistency. A test result can be checked for method, criteria and outcome. Only the second tells anyone whether the reviewer can do the job.

The committee's answer is already in the draft. A NOTE to 5.8.5 refers to ISO/IEC 9868 for recommendations on manual review processes and operator competence. A NOTE is informative, and a recommendation reached through one carries no presumption of conformity. The referral records the problem without resolving it. (A separate comment, number 13, observed that the four independence mechanisms are access controls and say nothing about a reviewer's interest in the outcome. That is a different point.)

2.2 Review of assignments does not review the assignee

Article 26(2) of the Act places the duty to assign oversight to persons with the necessary competence, training and authority on the deployer, and Article 4 places AI literacy on both provider and deployer. A provider standard cannot reach the deployer's staff, and the committee may have left the question there deliberately, or for ISO/IEC 9868, or for deployer-side guidance that does not yet exist. I cannot tell from the text. What the text does show is that the one channel from provider to deployer on this question, the declaration under 5.6.1, carries no signal when the system beneath it changes.

The second instance is about time. 5.2.6 is titled *Review of oversight assignments*. It requires the provider to review the risk scenarios, the reaction timeframes and the selected measures when post-market analysis or design modification alters the identified risks or their required oversight measures. It does not require review of the assignment itself: whether the person assigned still meets the competencies the provider declared.

Those competencies are declared under 5.6.1, in the instructions for use. My fifteenth comment proposed that the declaration be sufficient for a deployer to determine, before assignment, whether an individual meets it.

Nothing in the draft states when that determination stops holding. The design modification that triggers review under 5.2.6 is exactly the kind of change that invalidates it, and no clause assigns anyone to notice.

The natural answer is that this belongs to the sibling standard, and the draft says so: 5.2.6 performs its review in accordance with prEN 18228, Clause 11, titled *Risk management review*. That clause reviews the execution of the risk management plan, and its reassessment loop (Figure B.1) takes information about the AI system as its input: post-market findings, design modifications, incidents. The competence of assigned personnel is not among those inputs. 18228 does have a competence requirement, at 4.3: persons performing risk management activities shall be competent on the basis of education, training, skills and experience. It establishes competence at assignment. I found no later provision in the document that returns to it.

For the provider's own staff that silence matters less than it looks. Their competence is a quality management matter under Article 17, and quality management practice (ISO 9001 and ISO/IEC 42001, clause 7.2 in each) requires competence to be determined, evaluated, recorded and kept current. For the deployer's designated person it matters entirely. The provider's quality system does not reach that person, this document does not govern the deployer, and 5.6.1 is the only thread between them: a declaration of competencies written by the provider and read by the deployer, with no provision on either side for what happens when the system beneath it changes. 5.2.6 is the clause where the provider's change meets the deployer's person, and the referral to Clause 11 does not catch it.

3 · Where the specification stops

The pattern is not confined to this draft. In September 2026 two proposals appeared for new human oversight functions at the frontier. Dario Amodei proposed independent evaluators embedded inside AI developers with employee-like access. A joint statement of 22 heads of government and ministers, published by the Dutch government on 22 September and signed among others by the Dutch Prime Minister and the President of the European Commission, called for qualified evaluators to be granted sufficient access to assess risks. In their published text, both are specific about what those people would be given: access, independence, the right to report. Neither states what they must be able to do, or how anyone would know it still holds. The instinct is understandable; access, independence and reporting rights are the things one knows how to write down, and a standards committee working under a Commission mandate wrote down the same things.

The drafters were careful about people, in the register available to them. Actions can be specified because actions are observable, loggable and testable against a recorded criterion. The draft does all of that consistently, and most of the comments I wrote on those clauses tightened criteria already present in outline.

Capacities are different. To require that a person be able to detect an error, you need a method for finding out whether they can. Methods exist. Signal detection measures for error detection, vigilance paradigms, human reliability assessment in functional safety practice, recurrent proficiency checks in aviation and operator requalification in nuclear power are standardised and referenceable. None has been adapted to the review of AI output in a form this committee could cite, and the draft

cites none. Where it does borrow, at 5.3.4.5, it takes the metric names, time to intervention, misinterpretation rate and cognitive workload, and leaves the thresholds as a recommendation. The vocabulary confirms the position. *Automation bias* appears thirteen times, and the draft handles it in the right register, as a condition to be engineered through design measures at 5.5.1. *Reassessment*, *recertification*, *refresher* and *proficiency* appear zero times. *Competence* appears three times. *Competent* does not appear at all, in a standard for human oversight. Counts are from a whole-text search of the enquiry PDF, annexes and notes included, case-insensitive, on the stems shown; the search log is held by the author.

The finding is visible in the document's own structure. The draft specifies to the limit of what can currently be specified, and the limit falls between what a person does and what a person can do.

4 · What would close it

Two sentences, one for each document, would turn the declaration into something a deployer can act on when the system changes.

My sixteenth comment proposed one sentence for 5.2.6: where a design modification or post-market analysis requiring review under that subclause would invalidate a prior determination that a designated person meets the declared competencies, the provider shall identify that change in the instructions for use.

That clarifies a duty the provider already has under Article 13, to supply instructions for use, and under Article 43(4), to undergo a new conformity assessment after a substantial modification. It does not require the provider to assess the deployer's staff, which a provider standard cannot do. It requires the provider to say which of its own changes bear on a declaration the provider wrote.

A provider's counsel would object that "invalidate a prior determination" asks the provider to reason about a determination it never made and cannot see. A narrower drafting answers that: the provider shall state whether a modification changes the competencies declared under 5.6.1. Yes or no, against its own document. A deployer who sees the declaration change knows to look again at its people. The narrower form does the same work, and if the working group prefers it, the comment is served.

For prEN 18228 the equivalent is one sentence at 4.3: "Competence established under this subclause shall be confirmed at the intervals stated in the risk management plan and after any change that triggers review under Clause 11, and the outcome shall be recorded in the risk management file." That mirrors the document's own pattern for everything else it governs and assumes no method the document does not already assume.

Neither proposal solves the measurement problem. Each records a condition rather than measuring a capability. That is weaker than what the situation eventually needs and stronger than what either document currently has.

The measurement side is a deployer matter and sits outside both drafts. A record format for it is published separately: a competence matrix with the criterion and the date it was fixed, a round register completed before the

round runs, and a small function track for oversight functions too small to separate who builds a test from who takes it. (supervisory-capacity-assessment, CC BY 4.0; see Instrument under Sources.)

5 · Limitations

One reader, one enquiry draft of each standard, sixteen comments. I had no access to the committee's reasoning or to other commenters' submissions, and the drafts may have changed since the enquiry closed. The claims about the two September proposals rest on their published text as of 12 October 2026.

Sources and status

Terms. The provider places the AI system on the market. The deployer uses it. The designated person is the natural person to whom oversight of the system is assigned, in practice a member of the deployer's staff. The instructions for use are the provider's document to the deployer. The technical file is the provider's evidence for conformity assessment; Clause 6 of the draft lists what goes in it. Conformity with a harmonised standard is presumed conformity with the requirement of the AI Act it covers, here Article 14, so a provider that follows the standard is treated as meeting the Article, and what the standard does not ask for is unlikely to be checked.

Documents read. prEN 18229-3 (the human oversight part of the prEN 18229 series), public enquiry draft, enquiry closed 22 September 2026. prEN 18228 (AI risk management), public enquiry draft, obtained during the enquiry. Both drafted under standardisation request M/613 in support of Regulation (EU) 2024/1689.

Comments. Sixteen comments written during the enquiry, the last (comment 16, on 5.2.6) on 22 September 2026. Comments 9 (5.3.4.4), 13 (5.8.5), 15 (5.6.1) and 16 (5.2.6) are the ones this essay draws on directly and are reproduced in the Appendix. The full set is held by the author.

Proposals referred to in Section 3. Dario Amodei, "We Must Pace the Frontier", 12 September 2026, darioamodei.com. "A Call for Control of Frontier AI Models", diplomatic statement, 22 September 2026, published by the Netherlands Ministry of General Affairs, government.nl/documents/2026/09/22/a-call-for-control-of-frontier-ai-models; 22 signatories including Rob Jetten, Prime Minister of the Netherlands, and Ursula von der Leyen, President of the European Commission.

Methods referred to in Section 3. Signal detection: Green, D. M. and Swets, J. A. (1966), *Signal Detection Theory and Psychophysics*, Wiley. Vigilance: Mackworth, N. H. (1948), "The breakdown of visual vigilance during prolonged visual search", *Quarterly Journal of Experimental Psychology* 1(1), 6–21. Human reliability assessment in functional safety: IEC 61508 series; Kirwan, B. (1994), *A Guide to Practical Human Reliability Assessment*, Taylor & Francis. Recurrent proficiency checks in aviation: Commission Regulation (EU) No 1178/2011, Annex I (Part-FCL), proficiency checks for the revalidation of ratings. Operator requalification in nuclear power: US NRC, 10 CFR Part 55, Subpart F, requalification of licensed operators.

Underlying study. Turkina, D. (2026). The Unmonitored Dependency: Human Supervisory Capacity as an Assurance Target in Frontier AI Safety Frameworks. SSRN abstract 7248205; also at alphaxiv.org/abs/2609.frontier-ai-human-supervisory-capacity. Coding sheet and term-search log: doi.org/10.5281/zenodo.22075225 (CC BY 4.0).

Instrument. supervisory-capacity-assessment, release v1.0.0 (2026): record formats and a browser document builder for the deployer side. github.com/dturkina99/supervisory-capacity-assessment, CC BY 4.0.

Recurrence. The evidence comes in three tiers. Five frontier safety frameworks and two compliance filings were coded in the SSRN paper above and show no operative competence measurement. Two European standards drafts, read here, show a related pattern of a different kind: capacity tests without scope in one, competence at assignment with no return to it in the other. Two September 2026 proposals, read in their published text, specify evaluators by access and not by capacity.

Status of this document. Version 1.0, 12 October 2026, doi.org/10.5281/zenodo.23311590. Self-published as a record of observations made during the enquiry. Stages that follow: the working group resolves comments and prepares a Formal Vote draft, which is not public; the Commission's HAS consultants assess the draft against standardisation request M/613, including the Annex ZA mapping; the standard is published and cited in the Official Journal. If the published standard addresses either point, this document will be versioned to say so.

Licence. CC BY 4.0. The text of the draft standards is not reproduced here and remains under the licence of the issuing bodies.

Appendix · Four comments

Each comment is given with its clause reference, comment text and proposed change, as written.

Comment 9 · 5.3.4.4, the paragraph

Comment. The verification and its acceptance criteria are mandatory and the results are documented under 6.2 f), but three terms on which the requirement depends are not determined. "The key oversight tasks" has no antecedent; 5.3.4.5 a) refers to "a predefined oversight task" without stating who predefines it. "Designated persons" is defined at 3.4.3 by ascribed role and carries no further qualification here; the qualifier "representative" appears only in 5.3.4.5, which is a recommendation, and is not defined. The acceptance criteria have no metric set specified at requirement level: the metrics are listed in 5.3.4.5, where the provider is recommended, not required, to define acceptance thresholds for them, and the two subclauses use different terms ("criteria", "thresholds") without stating their relationship. 5.7.5 a) provides a model, testing against a response time defined in the risk management file. As drafted, the scope, population and pass condition of a mandatory verification are all set outside the requirement, and the result documented under 6.2 f) cannot be assessed for conformity.

Proposed change. Add at the end of 5.3.4.4: "The provider shall determine and record the oversight tasks subject to verification, the persons participating in the verification and the conditions under which it is performed. The acceptance criteria shall be defined over the metrics listed in 5.3.4.5. The tasks, participants and conditions recorded form part of the verification results documented in accordance with 6.2 f)."

Comment 13 · 5.8.5, item a)

Comment. 5.8.3 is titled "Manual reviewer independence mechanisms" and specifies four measures, all of which are access controls: individual credentials, prevention of a single account completing both reviews, blinding of the second reviewer, and separate recording of each review action. They establish that two distinct authenticated accounts performed the reviews. 5.8.5 specifies what the provider states about manual reviewers, and item a) lists three competencies, none concerning independence. No provision in the document states a condition on a reviewer's interest in the outcome of an identification or relationship to its subject; the NOTE to 5.8.5 refers to ISO/IEC 9868:2025 for recommendations on manual review processes and operator competence, but that reference is informative. The NOTE to 5.8.2.4 contemplates that manual reviewers may not reach the same assessment and leaves the resolution to a documented policy, while no provision states a condition under which a differing assessment is more likely to arise.

Proposed change. Add to 5.8.5, after item a) 3): "4) absence of any interest in the outcome of the identification, and of any relationship to its subject, that could affect the reviewer's assessment." Renumber the subsequent item accordingly.

Comment 15 · 5.6.1, third paragraph

Comment. The third paragraph provides that the competencies presumed of designated persons for output interpretation shall be those declared in the instructions for use, without stating what the declaration must contain. A declaration in general terms satisfies the requirement as drafted while conveying nothing a deployer can apply to a specific individual. Comparable specification duties elsewhere in the document state their required content: 5.4.1 b) 4) defines performance criteria as measurable indicators or success criteria; 5.3.2.1 lists the data elements the review interface shall contain; 5.4.3 d) requires a specification

sufficient for the deployer to act on without further information from the provider. The competencies presumed under this subclause are stated without any corresponding content requirement.

Proposed change. In 5.6.1, third paragraph, add: “The declaration shall be sufficient for the deployer to determine, before assignment, whether an individual meets the declared competencies.”

Comment 16 · 5.2.6, the paragraph

Comment. 5.2.6 is titled “Review of oversight assignments” and requires review of the risk scenarios, reaction timeframes and selected measures when post-market analysis or design modifications alter the identified risks or their required oversight measures. It does not require review of the assignment itself: whether the designated person still meets the competencies declared in accordance with 5.6.1. 5.6.1 makes those competencies a matter declared in the instructions for use, and a separate comment on 5.6.1 proposes that the declaration be sufficient for the deployer to determine, before assignment, whether an individual meets them. Neither clause states when that determination stops holding. The same design modifications that trigger review under 5.2.6 can invalidate it. 5.2.6 performs the review in accordance with prEN 18228:—, Clause 11. That review takes information about the AI system as its input; the competence of assigned personnel is not among its inputs and cannot trigger it. The gap therefore falls between the two documents rather than being resolved by the referral. This does not require the provider to assess the deployer’s personnel. It requires the provider to state which of its own modifications invalidate a prior determination, which only the provider is positioned to know.

Proposed change. Add to 5.2.6: “Where a design modification or post-market analysis requiring review under this subclause would invalidate a prior determination that a designated person meets the competencies declared in accordance with 5.6.1, the provider shall identify that change in the instructions for use.”